

Scalable and Memory-Efficient Clustering of Large-Scale Social Networks

Joyce Jiyoung Whang, Xin Sui, Inderjit S. Dhillon

Department of Computer Science

The University of Texas at Austin

IEEE International Conference on Data Mining (ICDM)

December 10 - 13, 2012

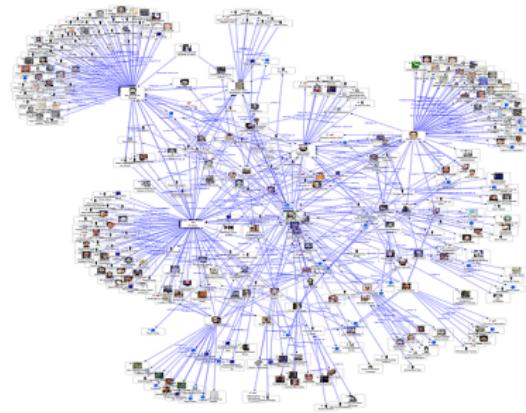
Contents

- Introduction
 - Problem Statement
 - Contributions
 - Preliminaries
- Multilevel Graph Clustering Algorithms
 - Multilevel Framework for Graph Clustering
 - Limitations of Multilevel Framework
- Proposed Algorithm: GEM
 - Graph Extraction
 - Clustering of Extracted Graph
 - Propagation and Refinement
- Parallel Algorithm: PGEM
- Experimental Results
- Conclusions

Introduction

Problem Statement

- Graph Clustering
 - Graph $G = (\mathcal{V}, \mathcal{E})$
 - k disjoint clusters $\mathcal{V}_1, \dots, \mathcal{V}_k$ such that $\mathcal{V} = \mathcal{V}_1 \cup \dots \cup \mathcal{V}_k$.
- Social Networks
 - Vertices: actors, edges: social interactions
 - Distinguishing properties
 - Power law degree distribution
 - Hierarchical structure



Contributions

- Multilevel Graph Clustering Algorithms
 - PMetis, KMetis, Graclus
 - ParMetis (Parallel implementation of KMetis)
 - Performance degradation
 - KMetis – 19 hours, more than 180 Gigabytes memory to cluster a Twitter graph (50 million vertices, one billion edges).
- **GEM (Graph Extraction + weighted kernel k -Means)**
 - Scalable & memory-efficient clustering algorithm
 - Comparable or better quality
 - Much faster and consumes much less memory
 - PGEM (Parallel implementation of GEM)
 - Higher quality of clusters
 - Much better scalability
 - GEM takes less than three hours on Twitter (40 Gigabytes memory).
 - PGEM takes less than three minutes on Twitter on 128 processes.



Preliminaries: Graph Clustering Objectives

- Kernighan-Lin objective
 - PMetis, KMetis
 - k equal-sized clusters

$$\min_{\mathcal{V}_1, \dots, \mathcal{V}_k} \sum_{i=1}^k \frac{\text{links}(\mathcal{V}_i, \mathcal{V} \setminus \mathcal{V}_i)}{|\mathcal{V}_i|} \text{ such that } |\mathcal{V}_i| = \frac{|\mathcal{V}|}{k}.$$

- Normalized cut objective
 - Graclus
 - Minimize cut relative to the degree of a cluster

$$\min_{\mathcal{V}_1, \dots, \mathcal{V}_k} \sum_{i=1}^k \frac{\text{links}(\mathcal{V}_i, \mathcal{V} \setminus \mathcal{V}_i)}{\text{degree}(\mathcal{V}_i)}.$$

Preliminaries: Weighted Kernel k -Means

- A general weighted kernel k -means objective is equivalent to a weighted graph clustering objective.
- Weighted kernel k -means
 - Objective

$$J = \sum_{c=1}^k \sum_{\mathbf{x}_i \in \pi_c} w_i \|\varphi(\mathbf{x}_i) - \mathbf{m}_c\|^2, \text{ where } \mathbf{m}_c = \frac{\sum_{\mathbf{x}_i \in \pi_c} w_i \varphi(\mathbf{x}_i)}{\sum_{\mathbf{x}_i \in \pi_c} w_i}.$$

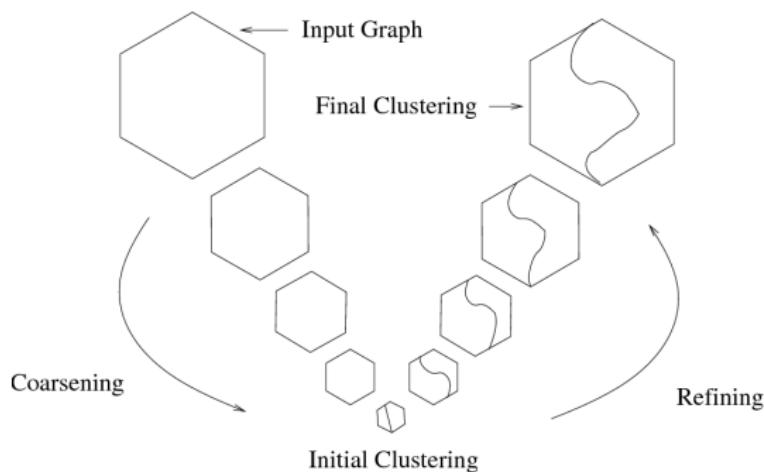
- Algorithm
 - Assigns each node to the closest cluster.
 - After all the nodes are considered, the centroids are updated.
 - Given the Kernel matrix K , where $K_{ij} = \varphi(\mathbf{x}_i) \cdot \varphi(\mathbf{x}_j)$,

$$\|\varphi(\mathbf{x}_i) - \mathbf{m}_c\|^2 = K_{ii} - \frac{2 \sum_{\mathbf{x}_j \in \pi_c} w_j K_{ij}}{\sum_{\mathbf{x}_j \in \pi_c} w_j} + \frac{\sum_{\mathbf{x}_j, \mathbf{x}_l \in \pi_c} w_j w_l K_{jl}}{(\sum_{\mathbf{x}_j \in \pi_c} w_j)^2}.$$

Multilevel Graph Clustering Algorithms

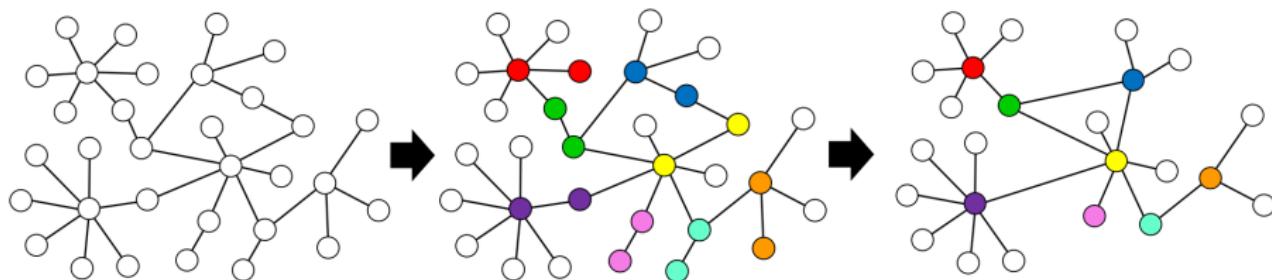
Multilevel Framework for Graph Clustering

- PMetis, KMetis, Graclus
- Multilevel Framework
 - Coarsening phase
 - Initial clustering phase
 - Refinement phase



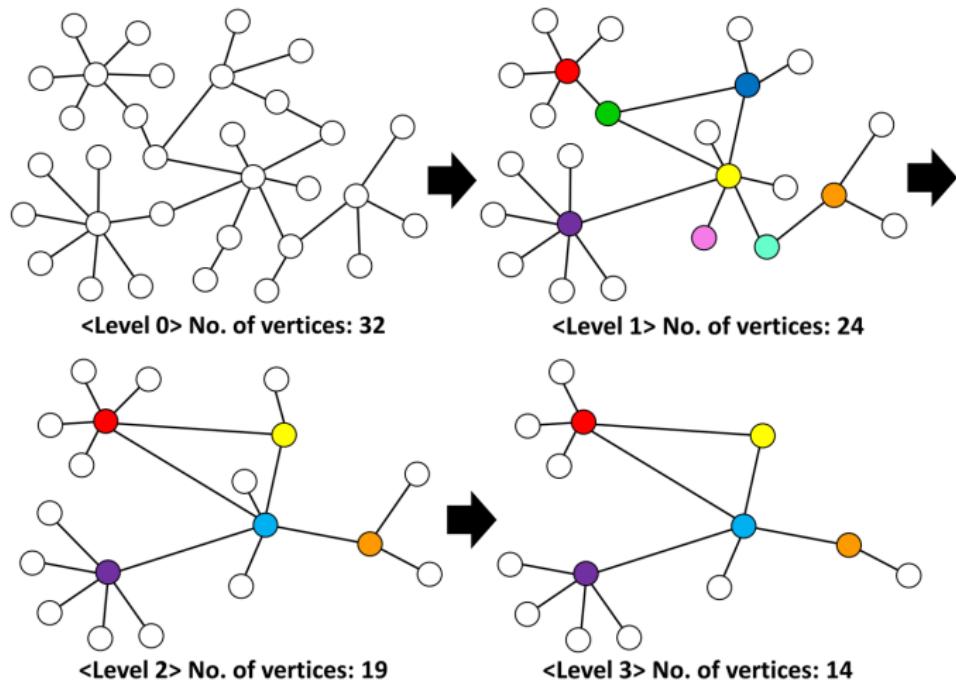
Problems with Coarsening Phase

- Coarsening of a scale-free network can lead to serious problems.
 - Most low degree vertices tend to be attached to high degree vertices.
 - Low degree vertices have little chance to be merged.
 - Example:
 - No. of vertices in the original graph: 32
 - No. of vertices in the coarsened graph: 24



Problems with Coarsening Phase

- Coarsening of a scale-free network
 - Transform the original graph into smaller graphs level by level

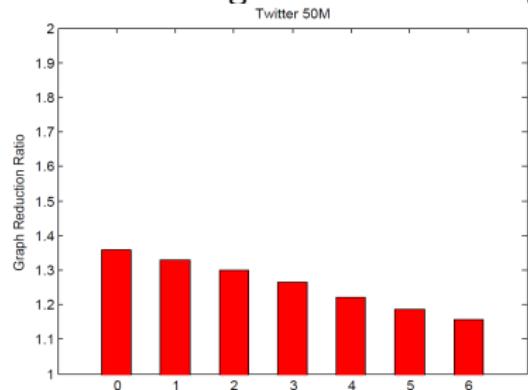


Limitations of Multilevel Framework

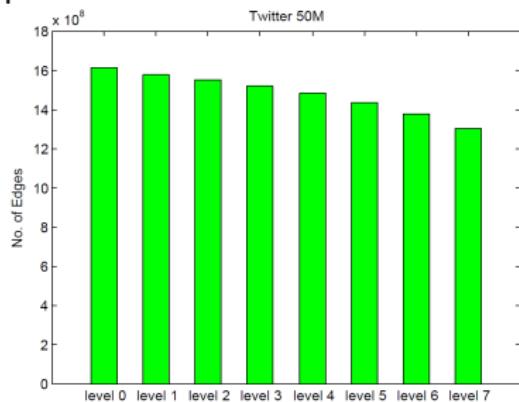
- Difficulties of Coarsening in Large Social Networks
 - In the coarsening from G_i to G_{i+1} ,

$$\text{Graph Reduction Ratio} = \frac{|\mathcal{V}_i|}{|\mathcal{V}_{i+1}|}.$$

- Ideally, graph reduction ratio would equal 2.
- The success of multilevel algorithms – high graph reduction ratio
- Power law degree distribution – graph reduction ratio becomes small.



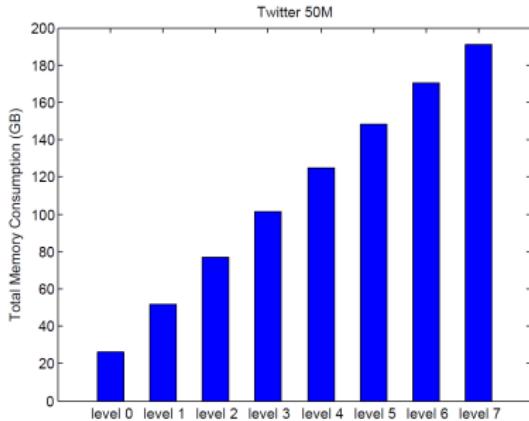
(a) Graph Reduction Ratio



(b) No. of Edges at Each Level

Limitations of Multilevel Framework

- Memory Consumption
 - Multilevel algorithms generate a series of graphs.
 - Total memory consumption increases rapidly during coarsening phase.

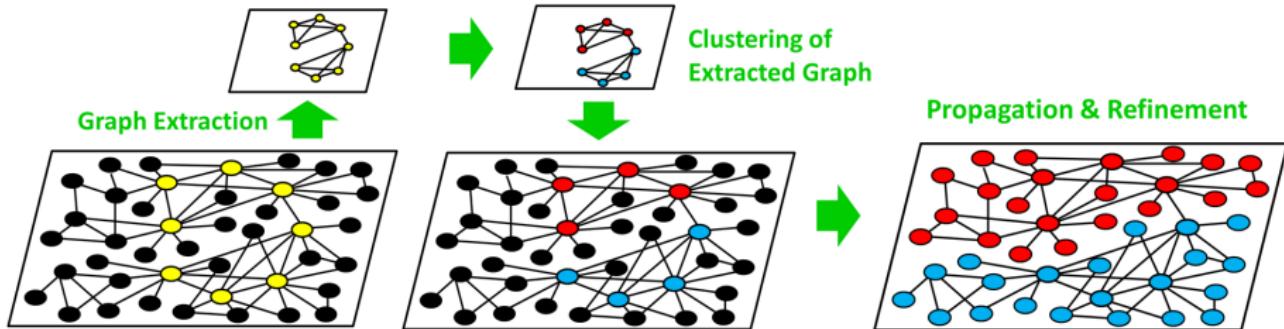


- Difficulties in Parallelization
 - Coarsening requires intensive communication between processes.

Proposed Algorithm: GEM

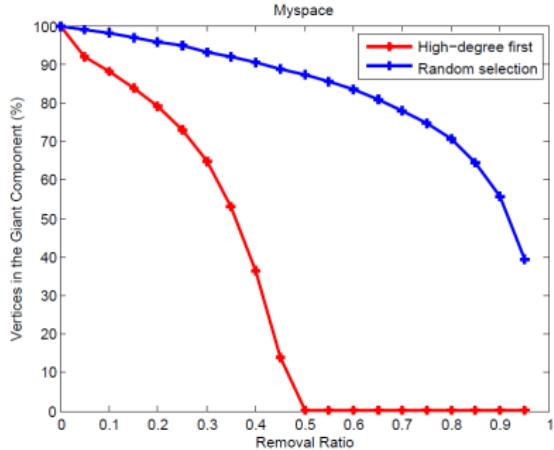
Overview of GEM

- Graph Extraction
- Clustering of Extracted Graph
- Propagation and Refinement



Graph Extraction

- Extract a skeleton of the original graph using high degree vertices.



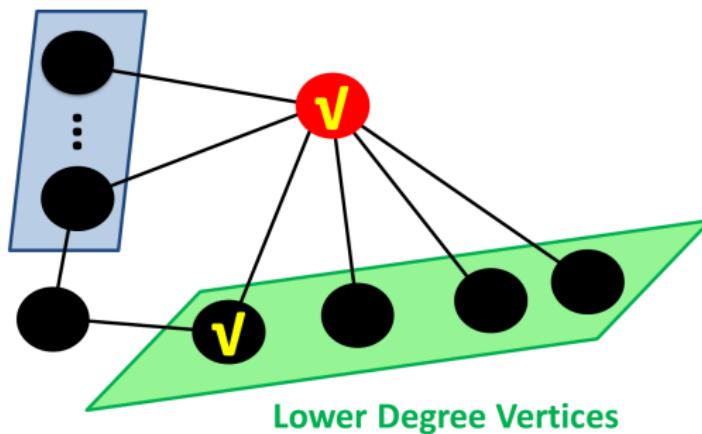
- High degree vertices
 - Tend to preserve the structure of a network
 - Popular and influential people

Clustering of Extracted Graph

- **Down-Path Walk Algorithm**

- Refer to a path $v_i \rightarrow v_j$ as a *down-path* if $d_i \geq d_j$.
- Follow a certain number of down-paths.
- Generate a seed by selecting the final vertex in the path.
- Mark the seed, and its neighbors.
- Repeat this procedure until we get k seeds in the graph.

Higher Degree Vertices

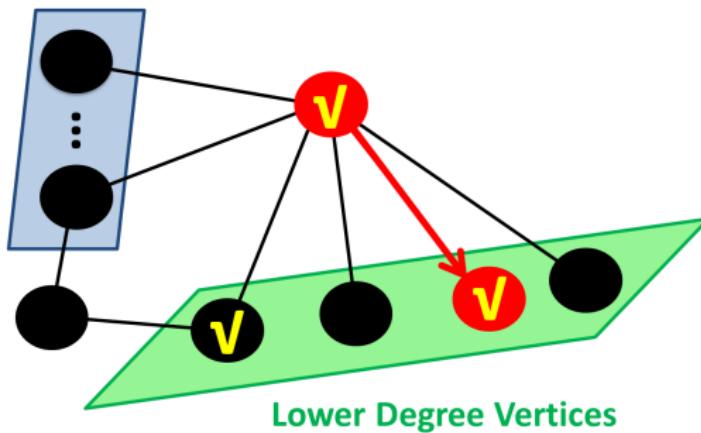


Clustering of Extracted Graph

- **Down-Path Walk Algorithm**

- Refer to a path $v_i \rightarrow v_j$ as a *down-path* if $d_i \geq d_j$.
- Follow a certain number of down-paths.
- Generate a seed by selecting the final vertex in the path.
- Mark the seed, and its neighbors.
- Repeat this procedure until we get k seeds in the graph.

Higher Degree Vertices



Clustering of Extracted Graph

- **Online Weighted Kernel k -means**

- Initialization of clusters

Input: \mathcal{V} : the vertex set, s_i ($i = 1, \dots, k$): seeds.

Output: \mathcal{V}_i ($i = 1, \dots, k$): initial clusters.

```
1: for each cluster  $\mathcal{V}_i$  do
2:    $\mathcal{V}_i = s_i$ .
3: end for
4: for each vertex  $\hat{v} \in \mathcal{V}$  do
5:   for each cluster  $\mathcal{V}_i$  do
6:      $\hat{\alpha} = \frac{\sigma}{\hat{d}} - \frac{2links(\hat{v}, \mathcal{V}_i)}{\hat{d} \cdot degree(\mathcal{V}_i)} + \frac{links(\mathcal{V}_i, \mathcal{V}_i)}{degree(\mathcal{V}_i)^2}$ .
7:      $\delta_i = \frac{\hat{d} \cdot degree(\mathcal{V}_i)}{degree(\mathcal{V}_i) + \hat{d}} \left\{ \hat{\alpha} + \frac{\sigma}{degree(\mathcal{V}_i)} \right\}$ .
8:   end for
9:   Find  $\mathcal{V}_p$  s.t.  $\delta_p \leq \delta_j$  for all  $j$  ( $j = 1, \dots, k$ ).
10:   $\mathcal{V}_p = \{\hat{v}\} \cup \mathcal{V}_p$ .
11: end for
```

Initialize clusters: each cluster
only contains the seed.

Assign each vertex to the cluster which allows
the least increase of the total objective.

Clustering of Extracted Graph

• Online Weighted Kernel k -means

- Graph clustering using online weighted kernel k-means

```
Input:  $\mathcal{V}$ : the vertex set,  $\mathcal{V}_i$  ( $i = 1, \dots, k$ ): initial clusters,  $\tau_{max}$ :  
maximum number of iterations.  
Output:  $\mathcal{V}_i^*$  ( $i = 1, \dots, k$ ): final clusters.  
1: Initialize  $\tau = 0$ .  
2: repeat  
3:   for each vertex  $\hat{v} \in \mathcal{V}$  do  
4:      $\mathcal{V}_p \leftarrow$  the current cluster of  $\hat{v}$ .  
5:     for each cluster  $\mathcal{V}_i$  do  
6:        $\hat{\alpha} = \frac{\sigma}{\hat{d}} - \frac{2links(\hat{v}, \mathcal{V}_i)}{\hat{d} \cdot degree(\mathcal{V}_i)} + \frac{links(\mathcal{V}_i, \mathcal{V}_i)}{degree(\mathcal{V}_i)^2}$ .  
7:       if  $\mathcal{V}_i = \mathcal{V}_p$  then  
8:          $\delta_i = \frac{\hat{d} \cdot degree(\mathcal{V}_i)}{degree(\mathcal{V}_i) - \hat{d}} \left\{ \hat{\alpha} - \frac{\sigma}{degree(\mathcal{V}_i)} \right\}$ .  
9:       else  
10:         $\delta_i = \frac{\hat{d} \cdot degree(\mathcal{V}_i)}{degree(\mathcal{V}_i) + \hat{d}} \left\{ \hat{\alpha} + \frac{\sigma}{degree(\mathcal{V}_i)} \right\}$ .  
11:       end if  
12:     end for  
13:     Find  $\mathcal{V}_q$  s.t.  $\delta_q \leq \delta_j$  for all  $j$  ( $j = 1, \dots, k$ ).  
14:     if  $\mathcal{V}_p \neq \mathcal{V}_q$  then  
15:        $\mathcal{V}_p = \mathcal{V}_p \setminus \{\hat{v}\}$ ,  $\mathcal{V}_q = \mathcal{V}_q \cup \{\hat{v}\}$ .  
16:       Update  $links(\mathcal{V}_p, \mathcal{V}_p)$ ,  $degree(\mathcal{V}_p)$ ,  
            $links(\mathcal{V}_q, \mathcal{V}_q)$ ,  $degree(\mathcal{V}_q)$ .  
17:     end if  
18:   end for  
19:    $\tau = \tau + 1$ .  
20: until not converged and  $\tau < \tau_{max}$   
21:  $\mathcal{V}_i^* = \mathcal{V}_i$  ( $i = 1, \dots, k$ ).
```

Compute the actual change in objective function for each cluster.

Assign the vertex to the cluster which allows the greatest decrease of the total objective change.

Propagation and Refinement

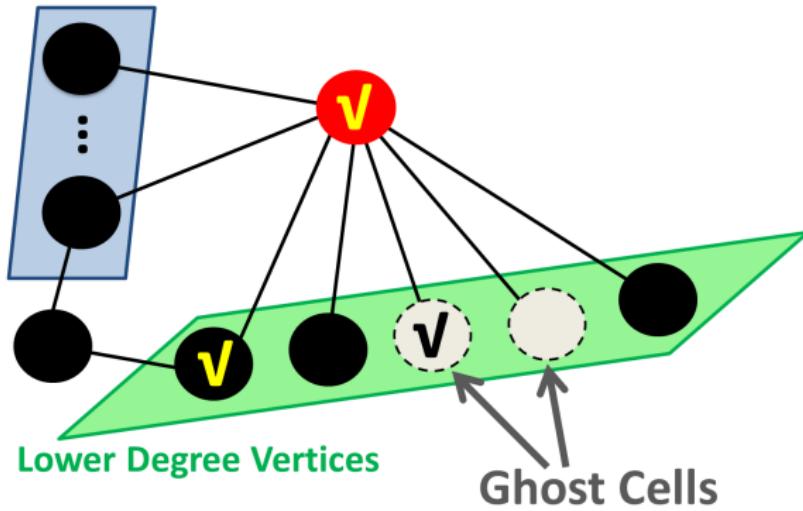
- Propagation
 - Propagate clustering of extracted graph to the entire original graph.
 - Visit vertices not in extracted graph in a breadth-first order (starting from vertices of extracted graph).
 - Arrive at initial clustering of the original graph (by online weighted kernel k -means).
- Refinement
 - Refine the clustering of the original graph.
 - Good initial clusters are achieved by the propagation step.
 - Refinement step efficiently improves the clustering result (by online weighted kernel k -means).

Parallel Algorithm: PGEM

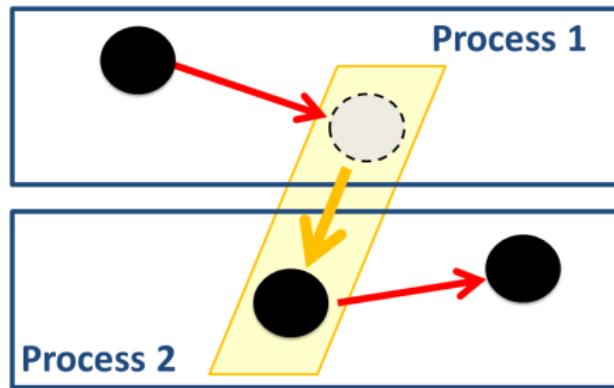
- GEM is easy to parallelize for large number of processes.
- Graph distribution across different processes
 - Each process: an equal-sized subset of vertices & their adjacency lists.
- Graph extraction
 - Each process scans its local vertices, and picks up high degree vertices.
 - The extracted graph is randomly distributed over all the processes.
- Seed selection phase
 - Seeds are generated in rounds.
 - Leader process decides the number of seeds each process will generate.
 - Proportional to the number of currently unmarked vertices in that process

- Parallel Down-Path Walk Algorithm
 - Neighbor vertices might be located in a different process.
 - Ghost Cells
 - For each remote neighbor vertex, a mirror vertex is maintained.
 - Buffer the information of its remote counterpart

Higher Degree Vertices



- Parallel Down-Path Walk Algorithm (continued)
 - Passing a walk to another process
 - Process 1 finds that the next vertex it will visit belongs to process 2.
 - Then process 1 stops this walk and notifies process 2 to continue it.



- Parallel Online Weighted Kernel k -means
 - Initialization of clusters
 - Ghost cells: accessing cluster information of remote vertices
 - Each process maintains a local copy of cluster centroids
 - Refinement
 - Visit local vertices in random order
 - Updating cluster centroids and ghost cells – relax the synchronization
- Propagation Phase
 - Similar to initialization of clusters in the extracted graph
- Refinement of the entire graph
 - Same strategy as clustering of the extracted graph

Experimental Results

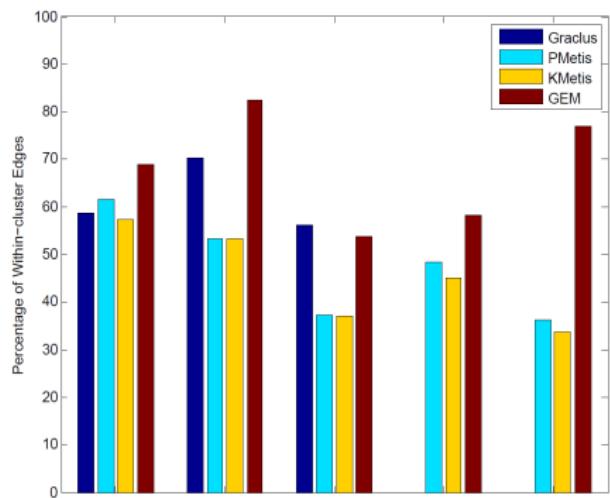
Experimental Setting & Dataset

- Sequential experiments
 - **GEM vs. PMetis, KMetis (Metis), Graclus**
 - Shared memory machine (AMD Opteron 2.6GHz CPU, 256GB memory)
- Parallel experiments
 - **PGEM vs. ParMetis**
 - Ranger at Texas Advanced Computing Center (TACC)
 - 3,936 machine nodes (4×4 -core AMD Opteron CPU, 32GB memory)
- Dataset

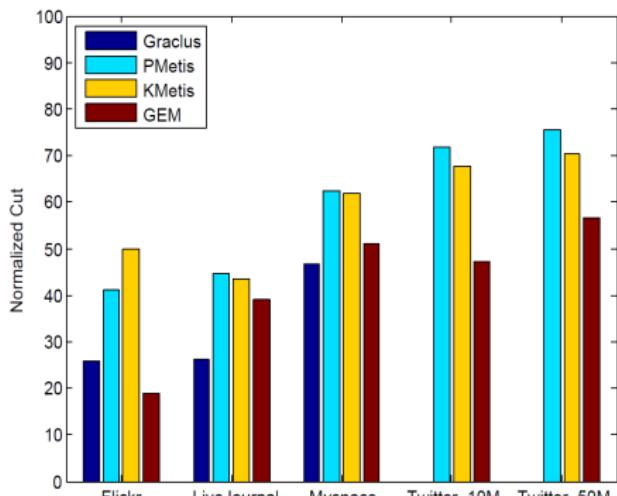
| Graph | No. of vertices | No. of edges |
|---------------|-----------------|---------------|
| Flickr | 1,994,422 | 21,445,057 |
| LiveJournal | 1,757,326 | 42,183,338 |
| Myspace | 2,086,141 | 45,459,079 |
| Twitter (10M) | 11,316,799 | 63,555,738 |
| Twitter (50M) | 51,161,011 | 1,613,892,592 |

Evaluation of GEM

- Quality of clusters
 - Higher percentage of within-cluster edges / lower normalized cut indicates better quality of clusters.



(a) Percentage of Within-cluster Edges

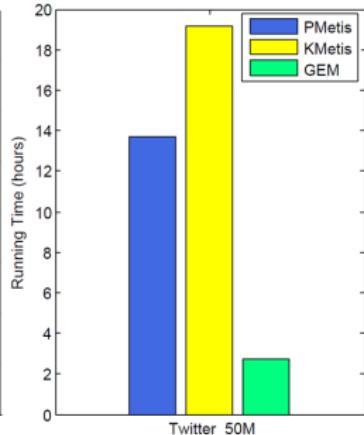
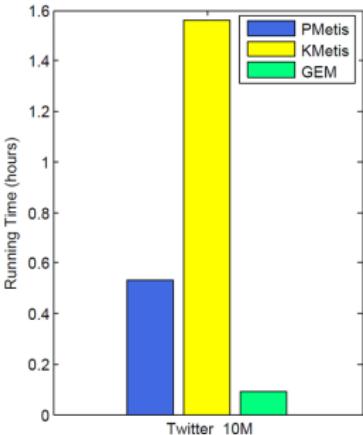
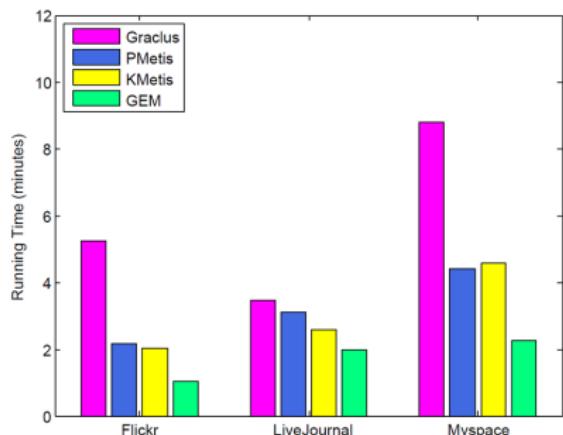


(b) Normalized Cut

Evaluation of GEM

- Running Time

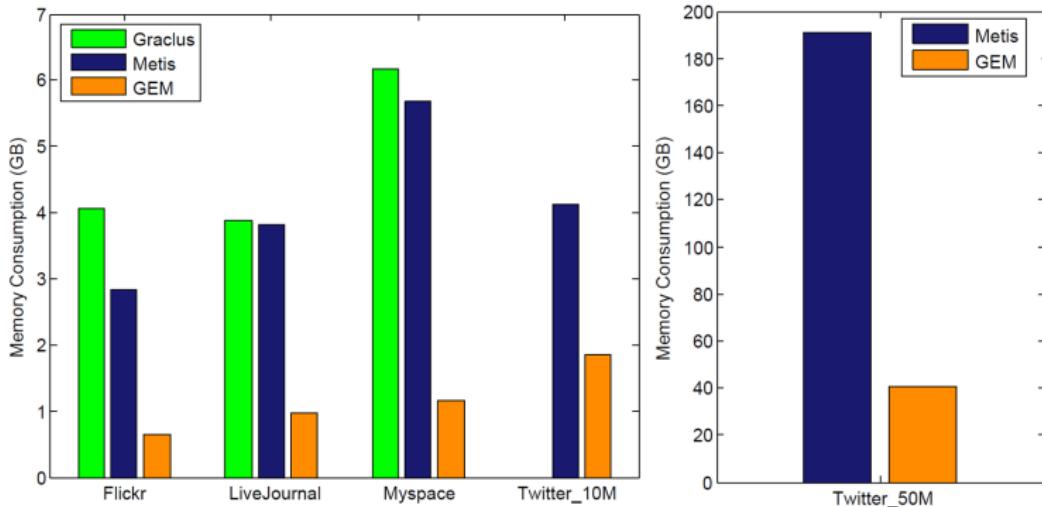
- GEM is the fastest algorithm across all the datasets.
- Twitter 10M: GEM (6 min.) vs. PMetis (30 min.) vs. KMetis (90 min.)
- Twitter 50M: GEM (3 hrs.) vs. PMetis (14 hrs.) vs. KMetis (19 hrs.)



Evaluation of GEM

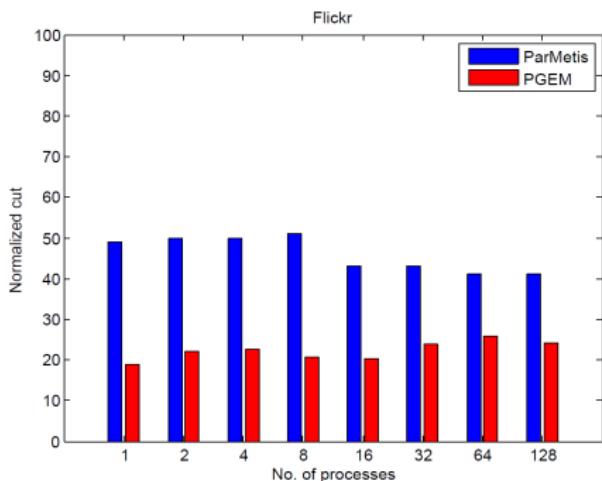
- Memory Consumption

- PMetis and KMetis use the same coarsening strategy (as in Metis).
- GEM directly extracts a subgraph from the original graph.
- Multilevel algorithms gradually reduce the graph size by multilevel coarsening.

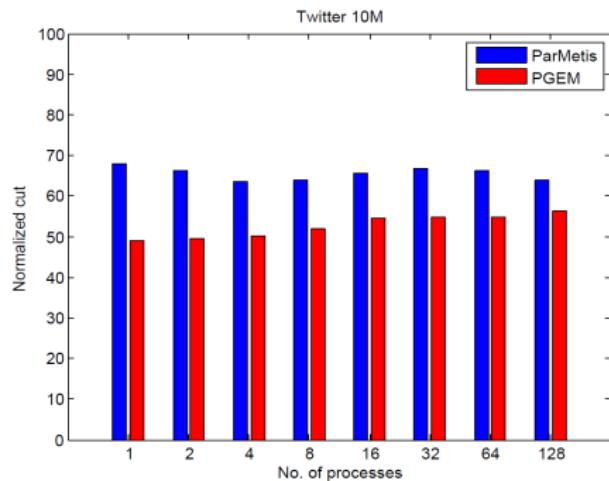


Evaluation of PGEM

- PGEM performs consistently better than ParMetis in terms of normalized cut, running time and speedup across all the datasets.
 - Quality of clusters



(a) Flickr normalized cut

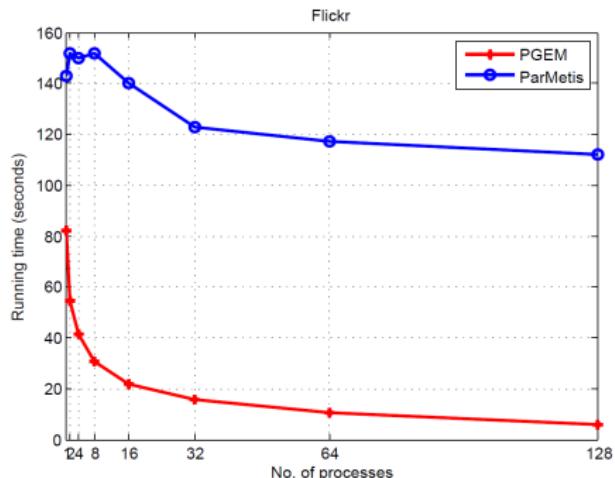


(b) Twitter 10M normalized cut

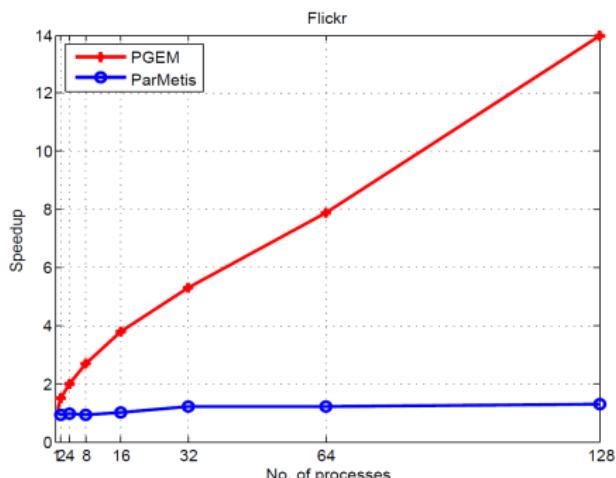
Evaluation of PGEM

- Running time & speedup on Flickr

$$\text{Speedup} = \frac{\text{Runtime of the program with one process}}{\text{Runtime with } p \text{ processes}}$$



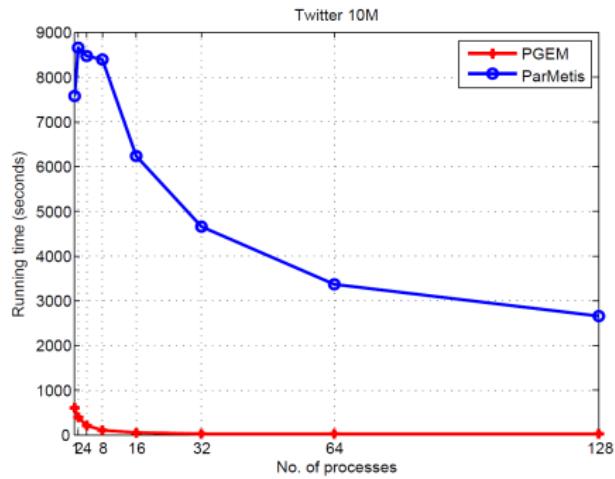
(a) Flickr running time



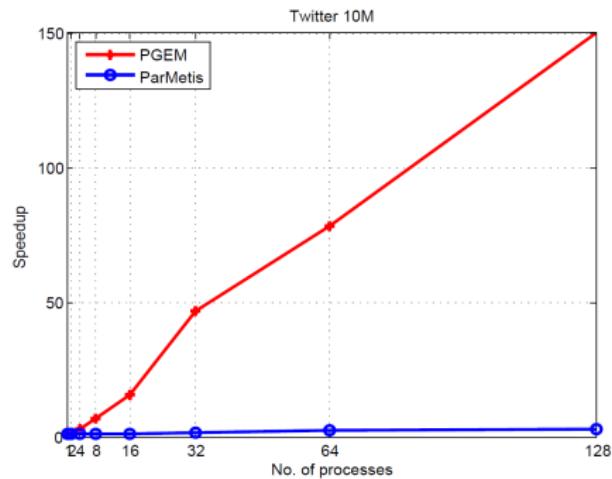
(b) Flickr speedup

Evaluation of PGEM

- Running time & speedup on Twitter 10M
 - PGEM achieves a super-linear speedup.
 - In all cases, the speedup of ParMetis is less than 10.
 - The multilevel scheme is hard to be scaled to large number of processes.



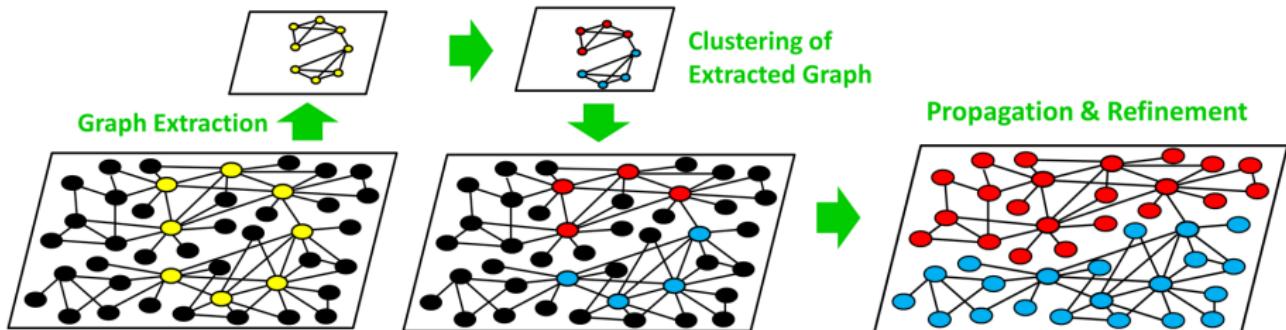
(a) Twitter 10M running time



(b) Twitter 10M speedup

Conclusions

GEM & PGEM



- GEM produces clusters of quality better than state-of-the-art clustering algorithms while it saves much time and memory.
- PGEM achieves significant scalability while producing high quality clusters.
- Future Research
 - Theoretical justification (can we theoretically show that extracted graph preserves structure of original graph).
 - Automatical detection of number of clusters.

References

- I. S. Dhillon, Y. Guan, and B. Kulis. Weighted graph cuts without eigenvectors: A multilevel approach. *IEEE Trans. on Pattern Analysis and Machine Intelligence*, vol. 29, no. 11, pp. 1944-1957, 2007.
- G. Karypis and V. Kumar. Multilevel k -way partitioning scheme for irregular graphs. *Journal of Parallel and Distributed Computing*, vol. 48, pp. 96129, 1998.